




c = 128, 45 min per cell, one vLLM pod per run. Dotted vertical line: end of the 10-minute warm-up. Hit rate is hits / queries in tokens over the trailing window, from the pod's own counters.