




4 vLLM pods behind one llm-d router (EPP); one 30-minute cell per arm and level (first 10 minutes warm-up); client timeout 1900 s; $x = c / 4$ for c in 16, 32, 48, 64, 96, 128, 192, 256, 338.