

Setup

model	Qwen3-Coder-30B-A3B-Instruct-FP8
engine	vLLM v0.28.0
replica	2 x H100 80GB, tensor parallel 2
KV cache	FP8, 48 KiB per token
GPU KV	2.24M tokens = 102 GiB
CPU tier	400 GiB = 8.74M tokens (copies GPU blocks, so it can hold about 8.74M tokens in total, not GPU + CPU)
workload	weka agentic coding traces; each turn resends the whole history; tool-call gaps capped at 10 s
router	one EPP per replica, 3 replicas in parallel

