

One vLLM replica per arm, offloading off (30 s rolling mean)

